


Generalized Linear Models

Prepared by: Joseph Bakarji

What if y is a label?




Given the data,
find a **function** h ,
that predicts y , given x

$$y = h(\mathbf{x})$$

$$y \in [0,1]$$


A **smooth** function that returns
probability of occurrence



What if y is a label?

$$y = h_{\theta}(x) \quad \& \quad y \in [0,1]$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-(\theta^{\top}x)}}$$



1. **Define a predictor:** the logistic function ✓
2. **Define a loss:** distance between function and data ?
3. **Optimize loss**
4. **Test model**


Logistic Regression

$$y = h_{\theta}(x)$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-(\theta^{\top} x)}} = g(\theta^{\top} x)$$

Linear predictor
negative log-likelihood or OLS

$$J(\theta) = \frac{1}{2} \sum_{i=1}^d \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$



Logistic predictor
Binary-cross entropy loss


$$\mathcal{L}(\theta) = \sum_{i=1}^n y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)}))$$

Gradient descent → Done!

Why not Least Squares?

$$y = h_{\theta}(x)$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-(\theta^{\top} x)}} = \sigma(\theta^{\top} x)$$



Probability of output given input

$$P(y = 1 \mid x; \theta) = h_{\theta}(x)$$

$$P(y = 0 \mid x; \theta) = 1 - h_{\theta}(x)$$



True label


$$p(y \mid x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$$

Likelihood!

Why not Least Squares?

$$y = h_{\theta}(x)$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-(\theta^{\top} x)}} = \sigma(\theta^{\top} x)$$



Probability of output given input

$$P(y = 1 \mid x; \theta) = \sigma(\theta^{\top} x)$$

$$P(y = 0 \mid x; \theta) = 1 - \sigma(\theta^{\top} x)$$



True label

$$p(y \mid x; \theta) = (\sigma(\theta^{\top} x))^y (1 - \sigma(\theta^{\top} x))^{1-y}$$

Likelihood!

Maximize **Log-likelihood**

$$\mathcal{L}(\theta) = \sum_{i=1}^n y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)}))$$

Update rule

while not converged:

$$\theta := \theta + \alpha \nabla_{\theta} \mathcal{L}(\theta)$$



Derive

Gradient Descent

for t = 1...T:

$$\theta := \theta - \alpha \sum_{i=1}^n \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x^{(i)}$$



Same as **linear regression**



Generalized Linear Models

Gaussian Distribution



Linear Regression

Bernoulli Distribution



Logistic Regression

Update rule

$$\theta := \theta - \alpha \sum_{i=1}^n \left(h_{\theta} (x^{(i)}) - y^{(i)} \right) x^{(i)}$$

Exponential Family

Family of distributions for which we can derive **the same update rule**

Assumption: $p(y | x; \theta)$ is an exponential family

Data ←

$$p(y; \eta) = b(y) \exp \{ \eta^\top y - a(\eta) \}$$

→ Parameters

- $b(y)$ is called the base measure (not depend on η)
- $a(\eta)$ is called the log partition function (not depend on y)
- $a(\eta)$, y and $b(y)$ are scalar. η and y have the same dimensions.

Example 1: Bernoulli Distribution -> Logistic Regression

Data \swarrow

$$p(y; \eta) = b(y) \exp \{ \eta^\top y - a(\eta) \}$$

\searrow **Natural** Parameters

Bernoulli Distribution

$$p(y; \phi) = \phi^y (1 - \phi)^{1-y} = \exp \left\{ y \log \frac{\phi}{1 - \phi} + \log(1 - \phi) \right\}$$

$\searrow \eta$ $\searrow a(\eta)$

**Show that term
is only a function of η**

Example 2: Gaussian Distribution -> Linear Regression

Data \leftarrow

$$p(y; \eta) = b(y) \exp \{ \eta^\top y - a(\eta) \}$$

\rightarrow Natural Parameters

Gaussian Distribution

$$p(y; \mu) = \frac{1}{\sqrt{2\pi}} \exp \left\{ -\frac{1}{2}(y - \mu)^2 \right\} = \underbrace{\frac{1}{\sqrt{2\pi}} e^{-y^2/2}}_{b(y)} \left\{ \underbrace{\mu y}_{\eta} - \underbrace{\frac{1}{2}\mu^2}_{a(\eta)} \right\}$$

Why do we care?

Data $\xleftarrow{\theta^\top x}$

$$p(y; \eta) = b(y) \exp \{ \eta^\top y - a(\eta) \}$$

$\xrightarrow{\text{Natural Parameters}}$

Inference is Easy:

$$E[y; \eta] = \frac{da(\eta)}{d\eta}$$

$$\text{Var}[y; \eta] = \frac{d^2 a(\eta)}{d\eta^2}$$

Learning is Easy:

Maximum Likelihood Estimation leads to **convex** problem in η

Generalized Linear Models

Assumption: $p(y | x; \theta)$ is an exponential family

Data Type → Probability Distribution

Binary → Bernoulli → **Logistic Regression**

Real → Gaussian → **Linear Regression**

Counts → Poisson

Positive Real → Gamma, Exponential

Distributions → Dirichlet

Generalized Linear Models

Assumption: $p(y | x; \theta)$ is an exponential family

The natural parameter is linear in the inputs

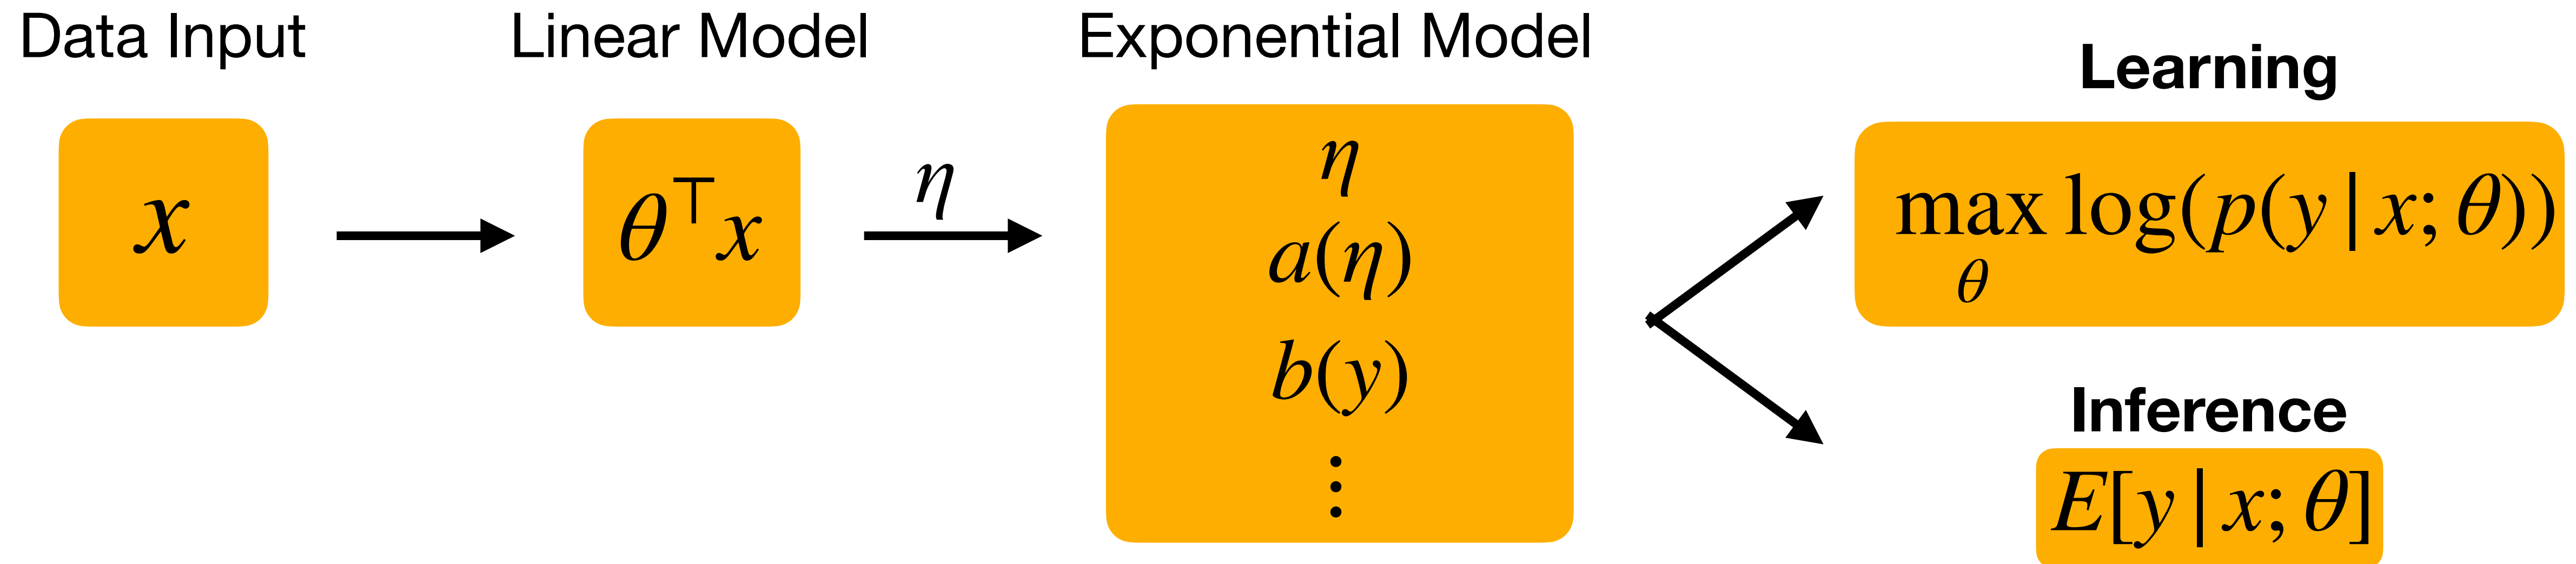
$$\eta = \theta^\top x$$

Predictor is a natural consequence

$$h_\theta(x) = E[y | x; \theta]$$

Generalized Linear Models

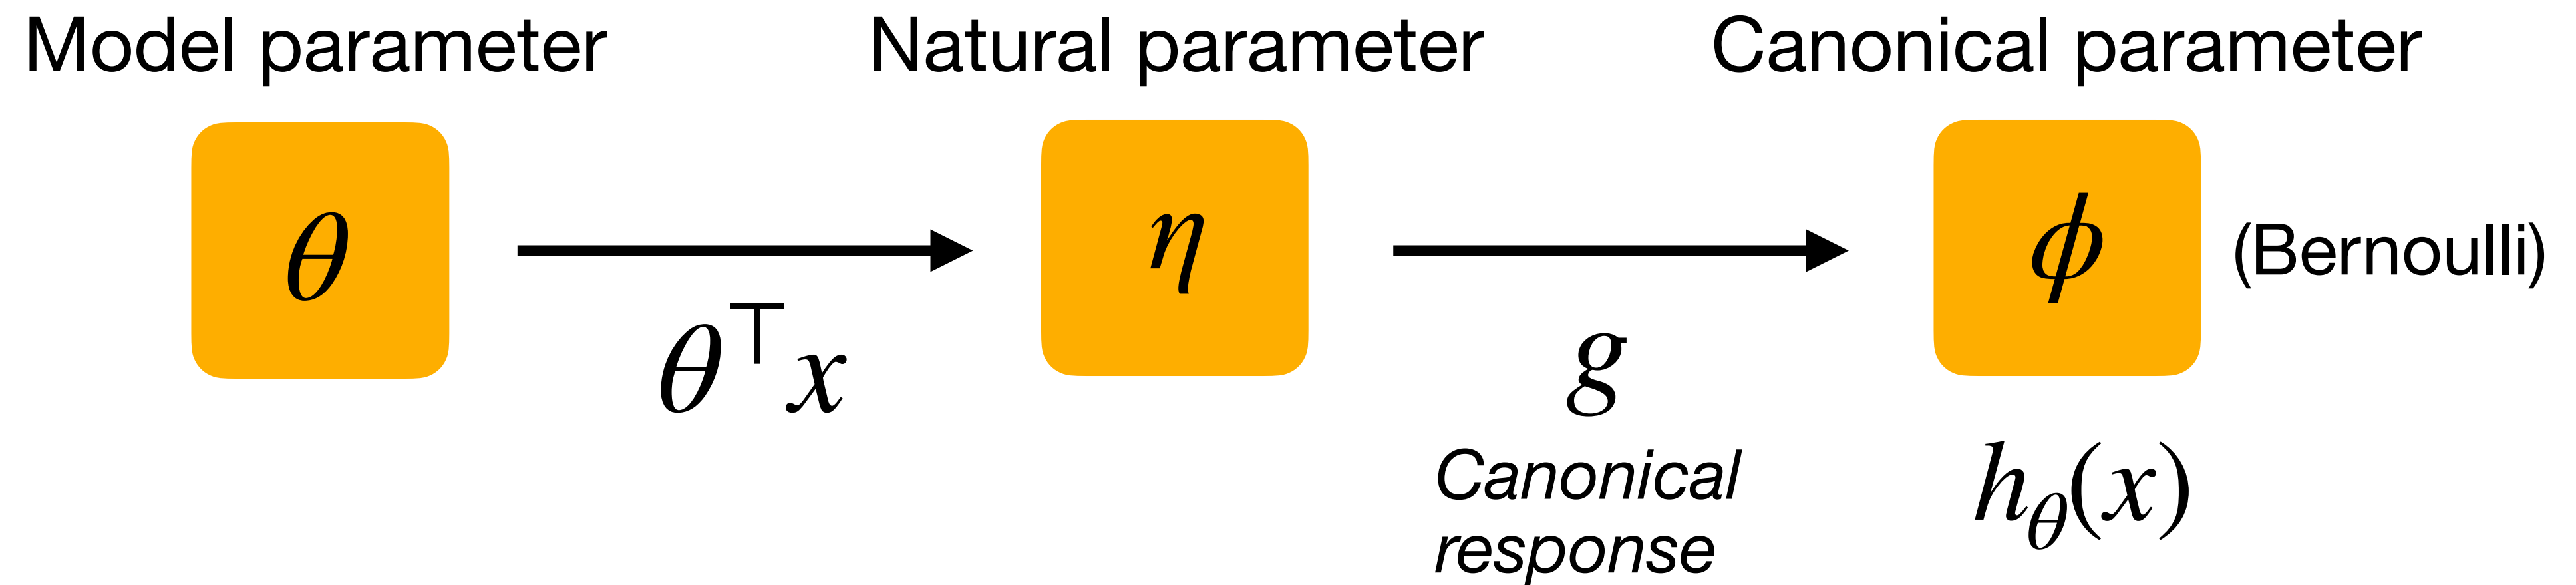
Assumption: $p(y | x; \theta)$ is an exponential family



Update Rule:

$$\theta := \theta - \alpha \sum_{i=1}^n \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right) x^{(i)}$$

Terminology




Logistic Regression:

$$h_\theta(x) = E[y | x; \theta]$$

$$\phi = \frac{1}{1 + e^{-\eta}} = \frac{1}{1 + e^{-\theta^\top x}}$$


Back to classification



$$y = h_{\theta}(x)$$

$$y \in [0,1]$$

What if we have more outputs?




Classification

$$x = [x_1, x_2]$$

x_1	x_2	y
-2	-1	●
3	1	★
2	3	●
1	-1	★
⋮		

★ 1

● 0



Logistic Regression

$$h_{\theta}(x) = \sigma(\theta^T x)$$



how confident?

score

$$\theta^T x$$


how correct?

margin

$$(\theta^T x)y$$

For $y \in [1, -1]$

Multiclass classification - Softmax



k discrete values for representing output


$$\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

car

$$\begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

plane

$$\begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

boat

$$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

horse

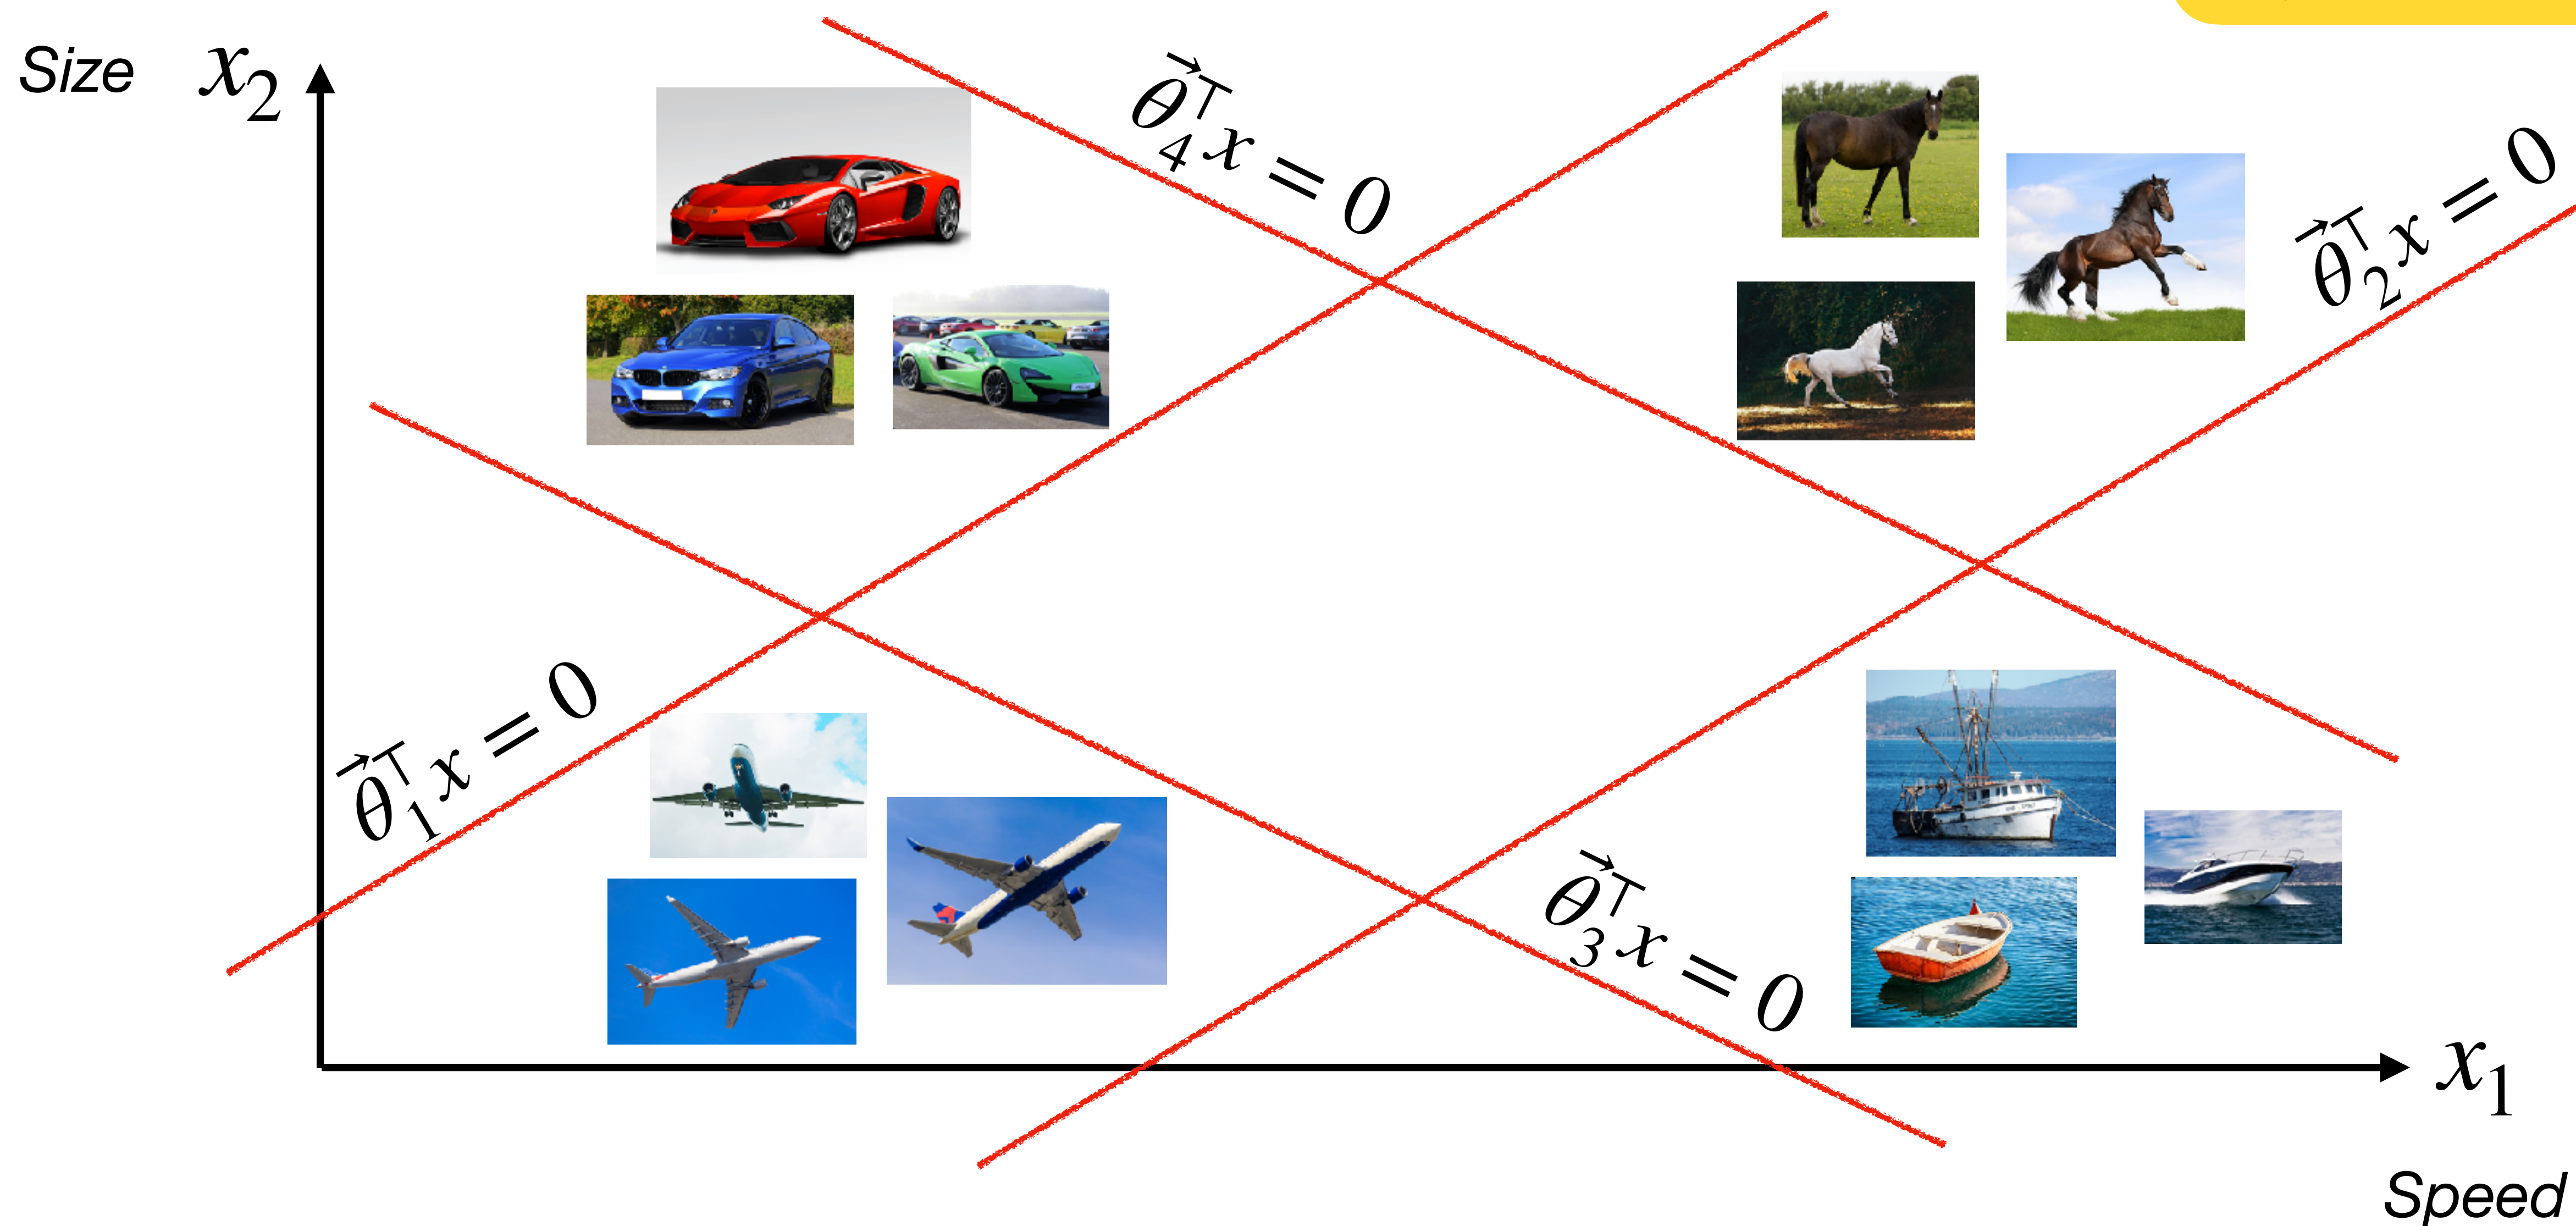
One-hot encoding

Multi-class classification - **Softmax**


WARNING!!!

Notation Alert

$\vec{\theta}_i$ is a vector



How to turn **scores** into **probabilities**?



WARNING!!!
Notation Alert
 $\vec{\theta}_i$ is a vector

Score

$$\begin{aligned}\vec{\theta}_1^T \hat{x} &= 3 \\ \vec{\theta}_2^T \hat{x} &= -0.3 \\ \vec{\theta}_3^T \hat{x} &= -0.8 \\ \vec{\theta}_4^T \hat{x} &= -22\end{aligned}$$

exp

Positive Measure


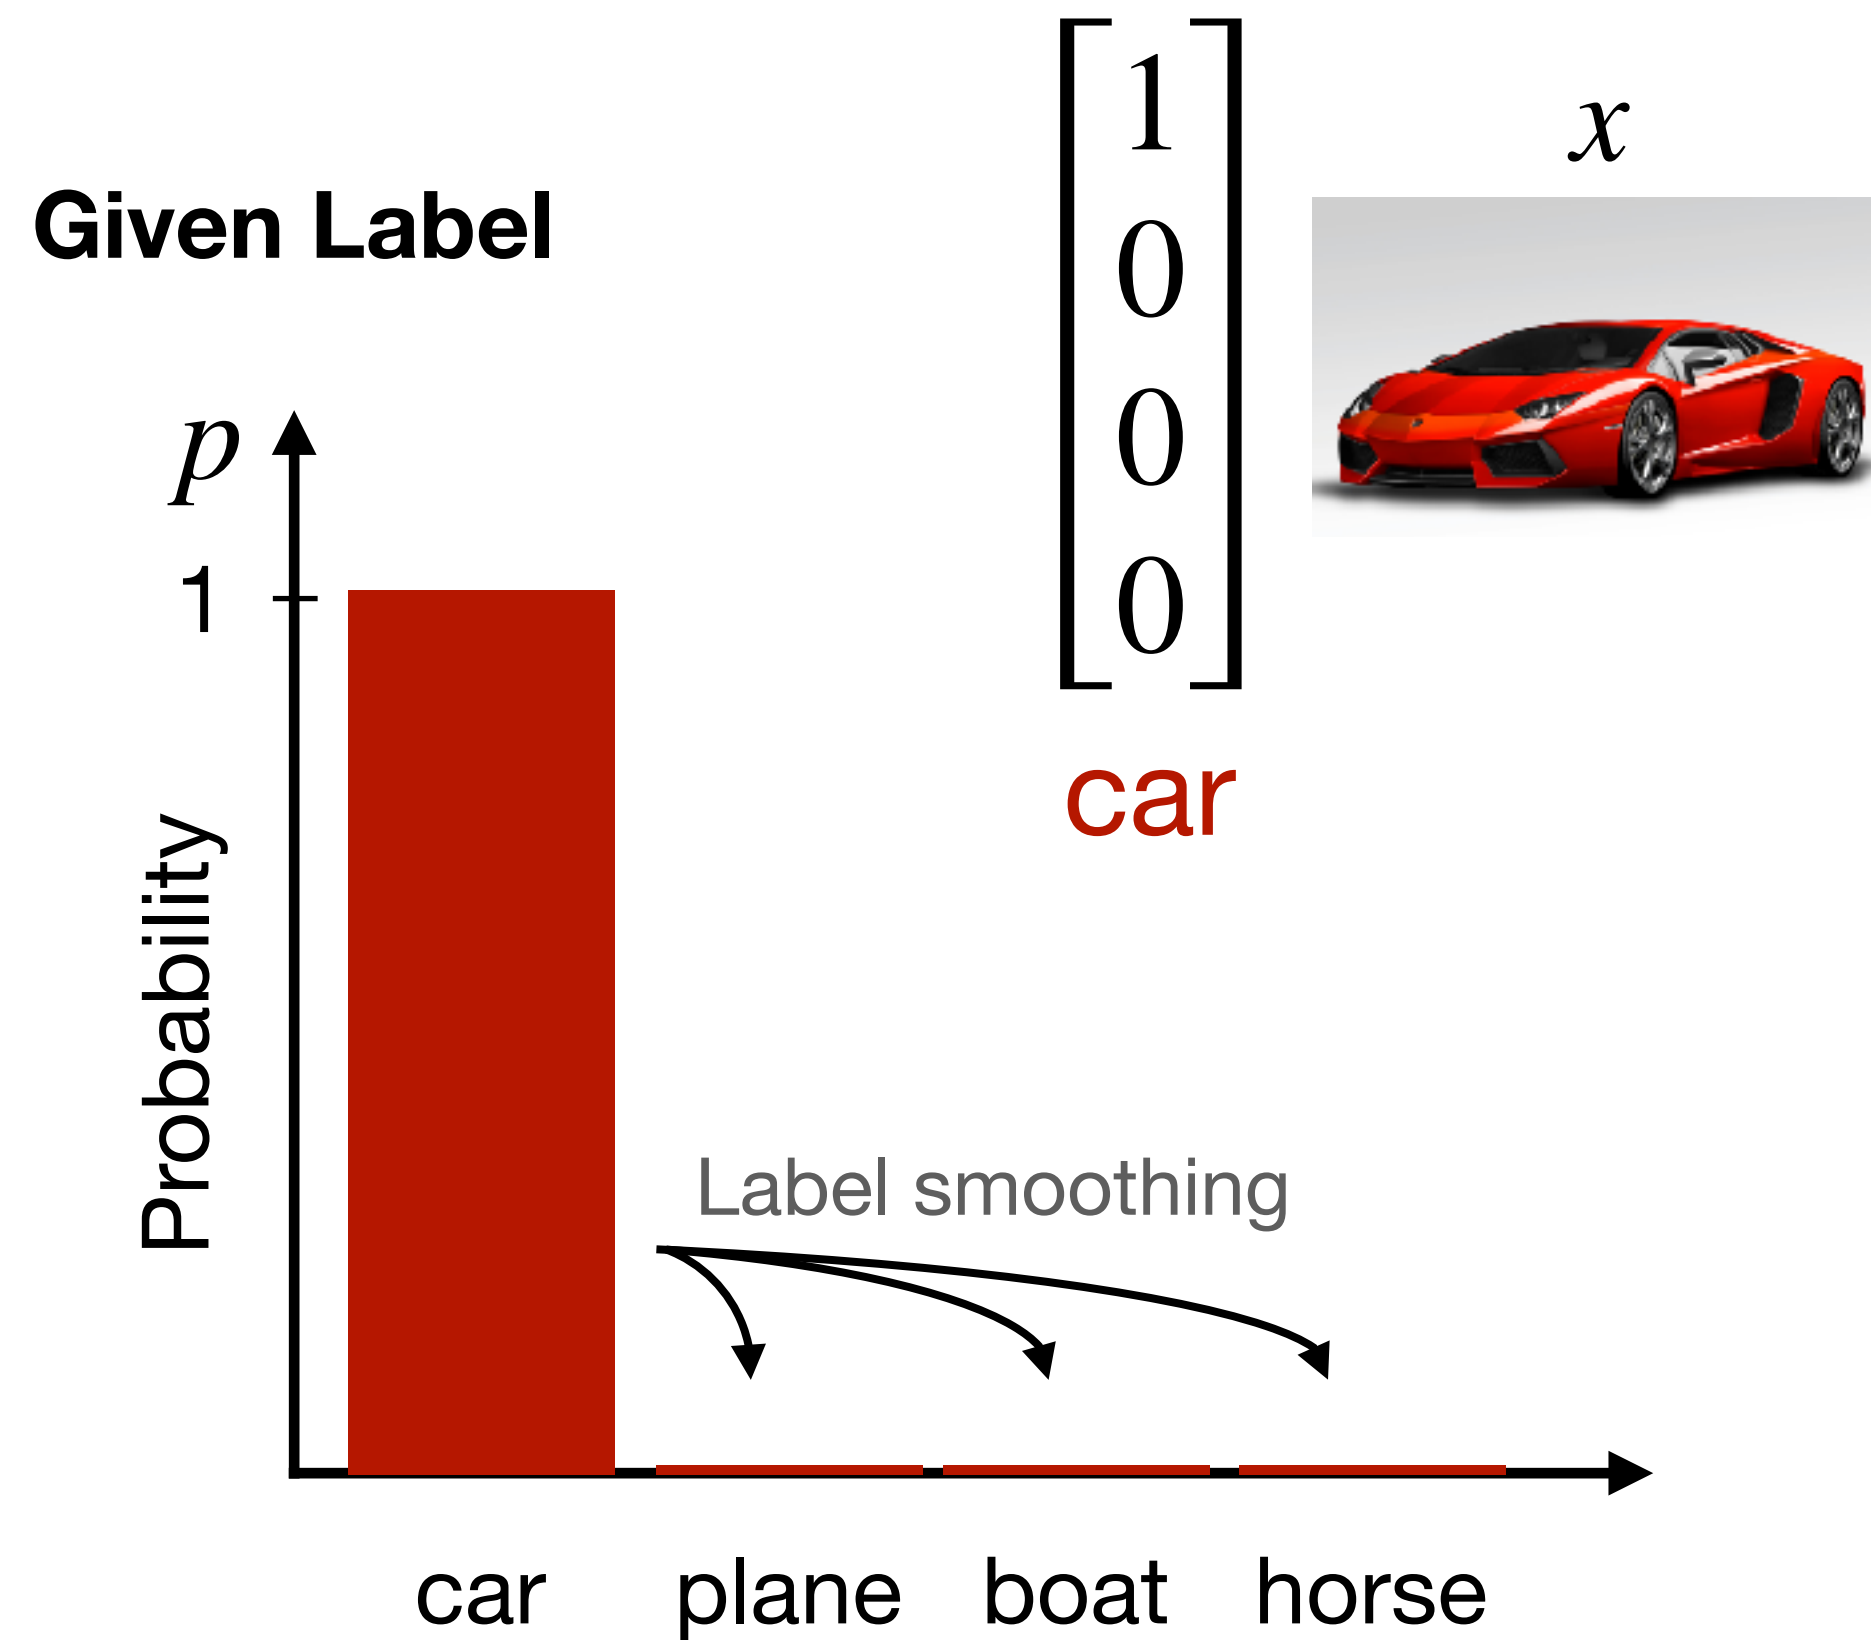
$$\begin{aligned}\exp(3) &= 20.1 \\ \exp(-0.3) &= 0.75 \\ \exp(-0.8) &= 0.2 \\ \exp(-22) &= 0.00..1\end{aligned}$$

Normalize

Softmax


$$\hat{p}(y = i | x; \theta) = \frac{\exp(\vec{\theta}_i^T x)}{\sum_{j=1}^k \exp(\vec{\theta}_j^T x)}$$

How do you train?



$$\begin{aligned} \min \text{CrossEntropy}(p, \hat{p}) &= - \sum_{i=1}^k p(y = i) \log (\hat{p}(y = i)) \\ &= - \log (\hat{p}(y = 1)) \end{aligned}$$

How do you train?



$$\text{CrossEntropy}(p, \hat{p}) = - \sum_{i=1}^k p(y = i) \log (\hat{p}(y = i))$$

Ground Truth

$$\text{Logit} = - \log (\hat{p}(y = 1))$$

$$= - \log \left(\frac{\exp(\vec{\theta}_i^\top x)}{\sum_{j=1}^k \exp(\vec{\theta}_j^\top x)} \right)$$

Train with Gradient Descent!